

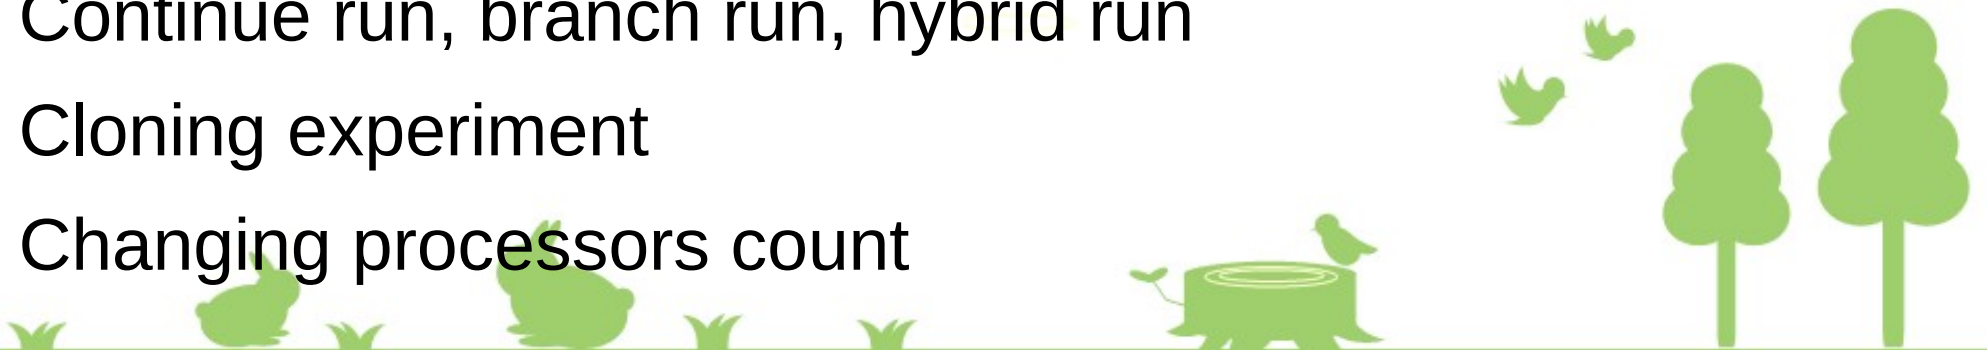
NorESM user workshop 2020

Alok K. Gupta, Yanchun He, Mats Bentsen



Agenda

- HPC infrastructures and NorESM workflow
- Downloading NorESM
- NorESM-2 structure
- Basic steps to execute model
- Few namelist variables to set up model and writing restart files
- user_nl_* files to modify namelist; SourceMods to debug source files
- Continue run, branch run, hybrid run
- Cloning experiment
- Changing processors count



HPC machines

- Fram

System Lenovo NeXtScale nx360

Number of Cores 32256

Number of nodes 1006

CPU type Intel E5-2683v4 2.1 GHz

Intel E7-4850v4 2.1 GHz (hugemem)

Per node 32 cores; 64 GiB memory

Island topology; 9216 cores per island

- We are using Intel compilers on both machine

- Betzy

BullSequana XH2000

Number of Cores 172032

Number of nodes 1344

CPU type AMD® Epyc™ 7742
2.25GHz

Per node 128 cores; 256 GiB memory

- Recently, in testing phase and have approx. 15-25% efficiency for most experiments of NorESM

Information taken from Sigma2 website :- <https://www.sigma2.no/>




Web services:

- Research Data Archive : <https://archive.norstore.no/>
- Diagnostic output : <http://ns2345k.web.sigma2.no/>
- ESGF node : <https://noresg.nird.sigma2.no/thredds/>
- NorESM git repository : <https://github.com/NorESMhub>
- NIRD toolkit : <https://apps.sigma2.no/nird>
- NorESM documentation : <https://noresm-docs.readthedocs.io/en/latest/>
- NorESM inputdata Server : <https://www.noresm.org/inputdata>



Recent technical development

- Distributed NorESM repository – Like as CESM- public
- Support for new system Betzy
- More documentations – NorESM user guide
- MICOM changes to BLOM – Bergen Layered Ocean Model
- user_nl_blom – consistent like as user_nl_cam ...
- Few bug fixes related to uninitialized variables
- Implementation of inputdata server:- CAM-OSLO and BLOM specific data can be downloaded automatically
- Implementation of CMIP6 compsets
- License



Workflow

Model git repository
cime, cam, blom, cice,
clm , rtm, ..

<https://github.com/NorESMhub>

Model Simulation
Mostly, shared folder


FRAM – Betzy – HPC Machines

Post-processing
Diagnostics,
CMORization

NIRD - National e-Infrastructure for Research Data

Published and archived:
CMIP6 , CMIP5 ...

Longterm storage



Downloading CMIP6 version of NorESM

- git clone – clone the remote repository locally

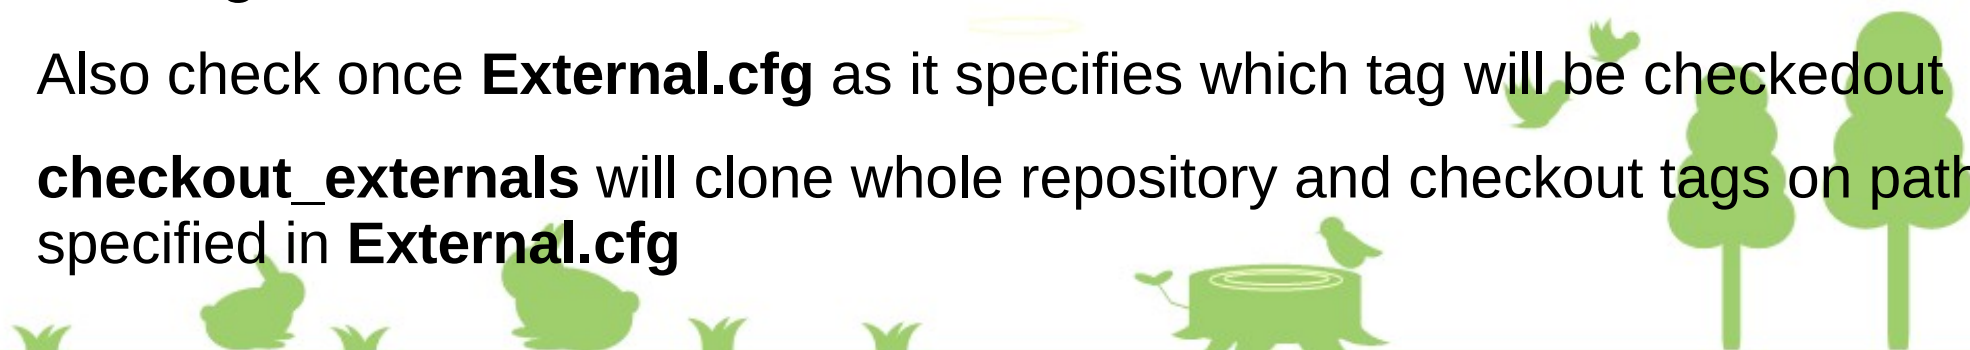
```
git clone https://github.com/NorESMhub/NorESM NorESM
```

```
cd NorESM
```

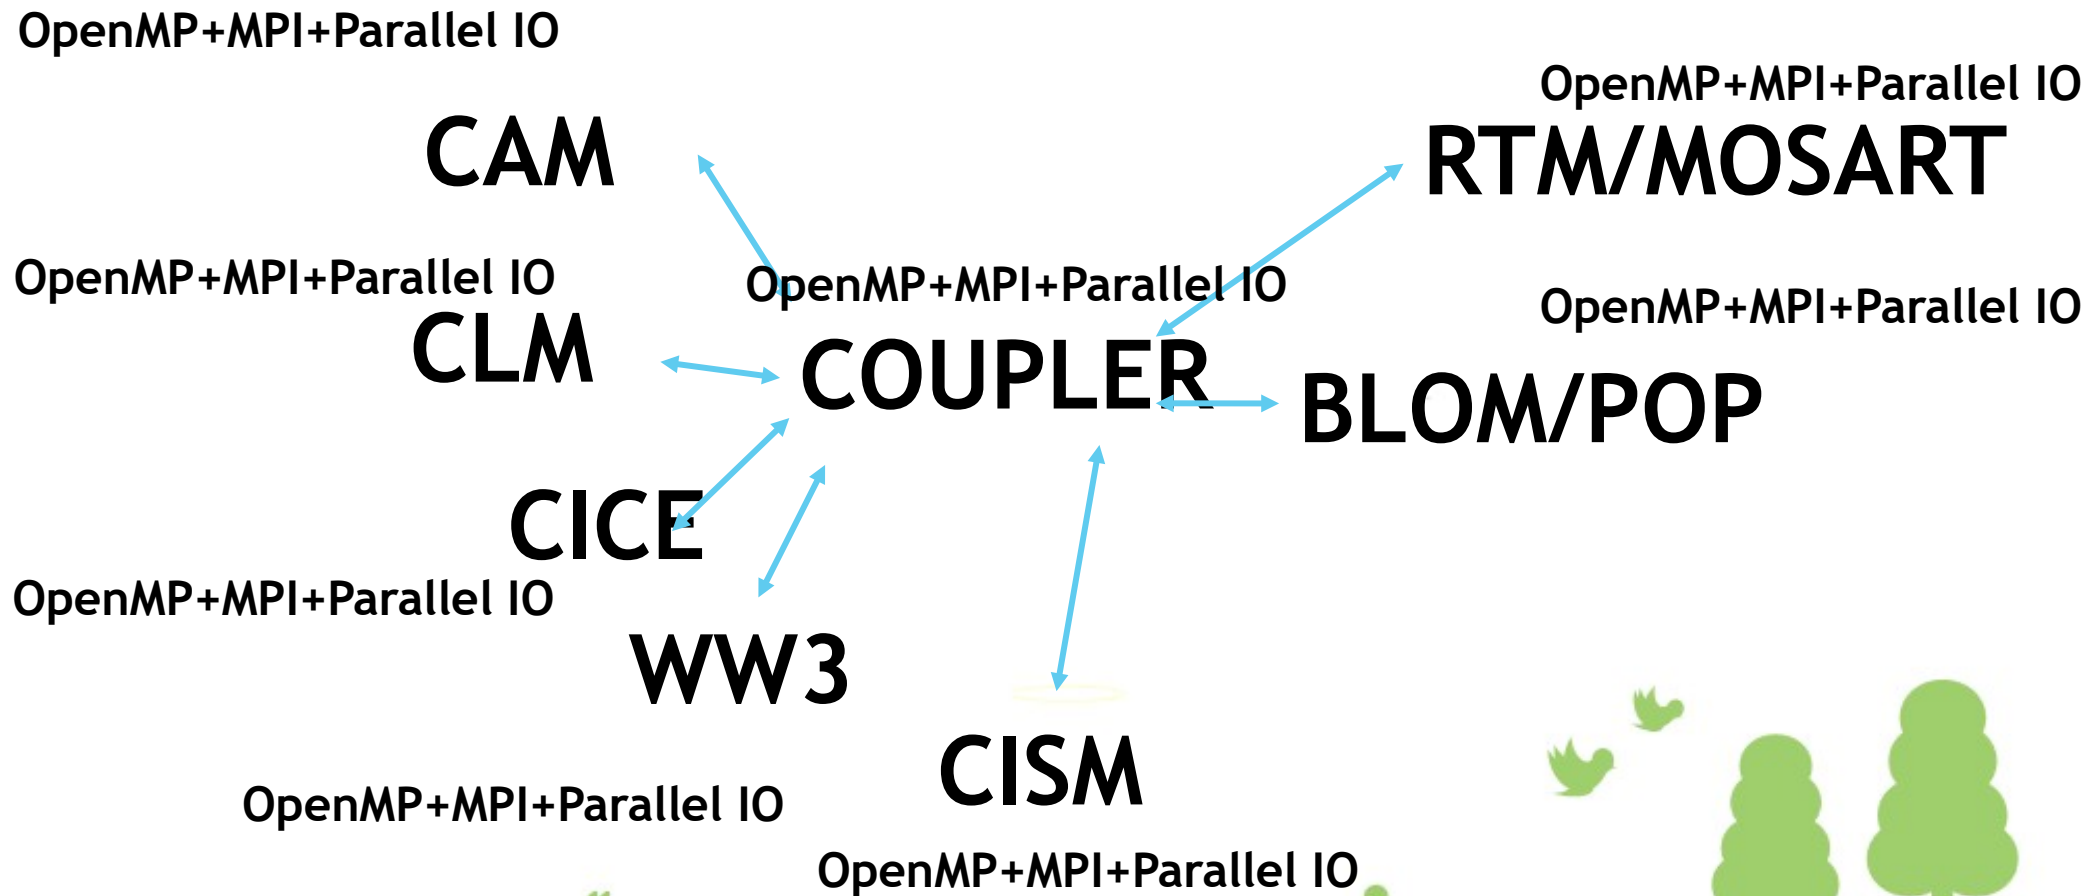
```
git tag    list all tags
```

```
git checkout tags/release-noresm2.0.2 -b noresm
```

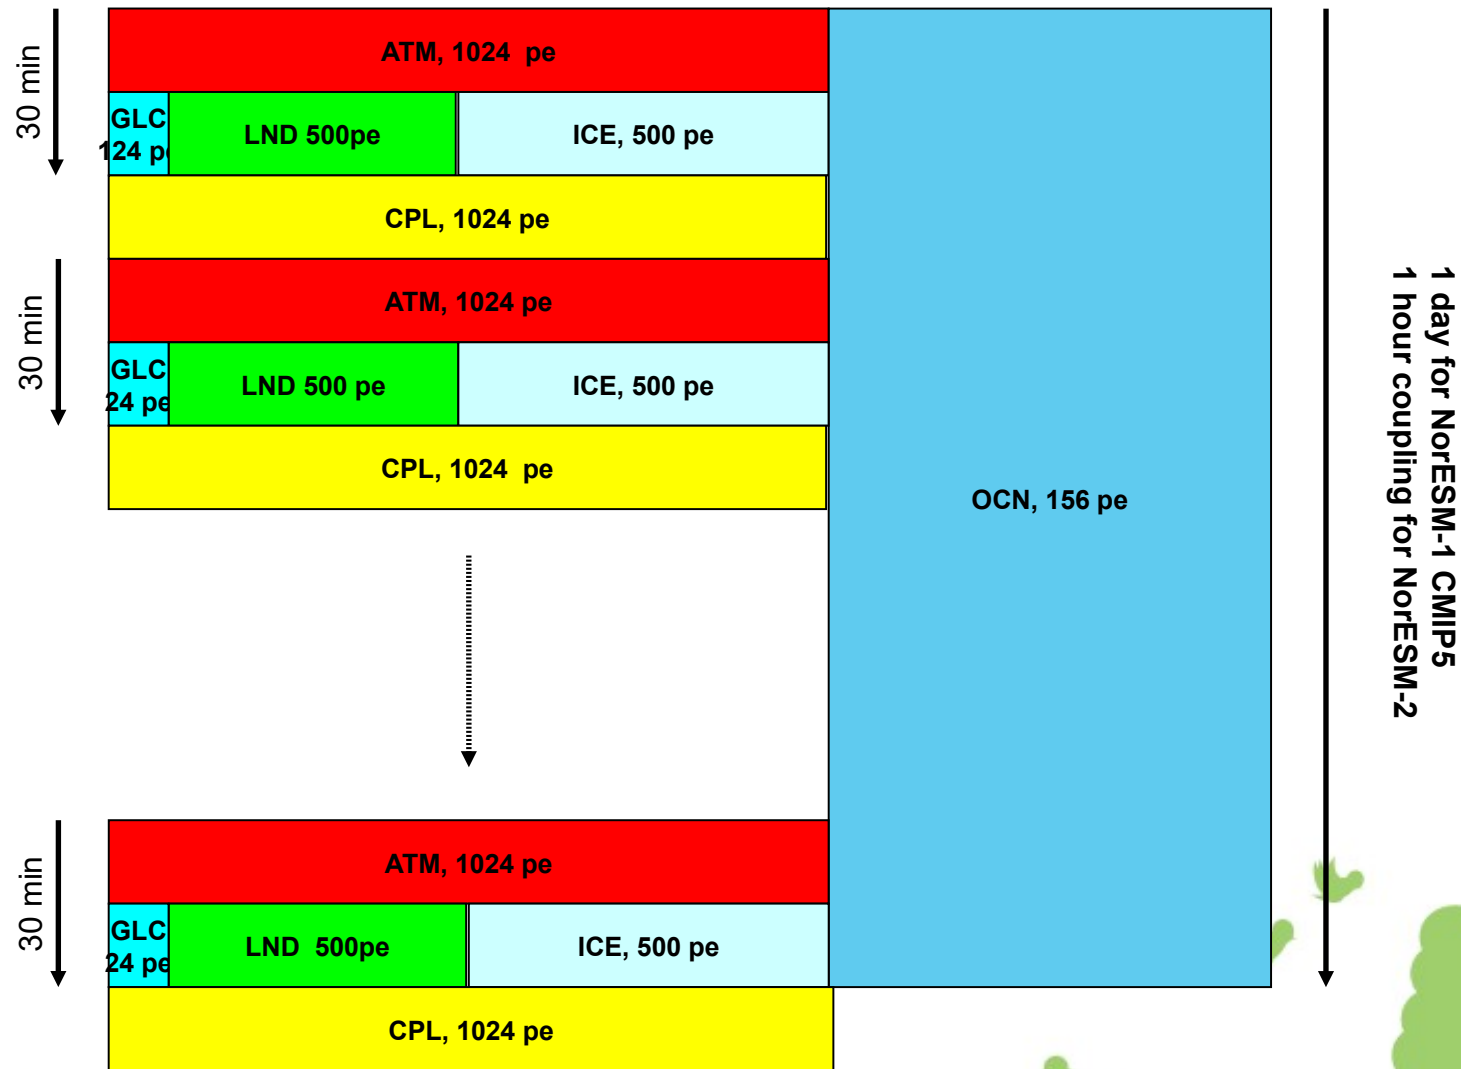
- checkout tag **release-noresm2.0.2** to **noresm**
- **release-noresm2.0.2** – version to reproduce CMIP6 results of NorESM2
./manageExternals/checkoutExternals
- Also check once **External.cfg** as it specifies which tag will be checkedout
- **checkoutExternals** will clone whole repository and checkout tags on path specified in **External.cfg**



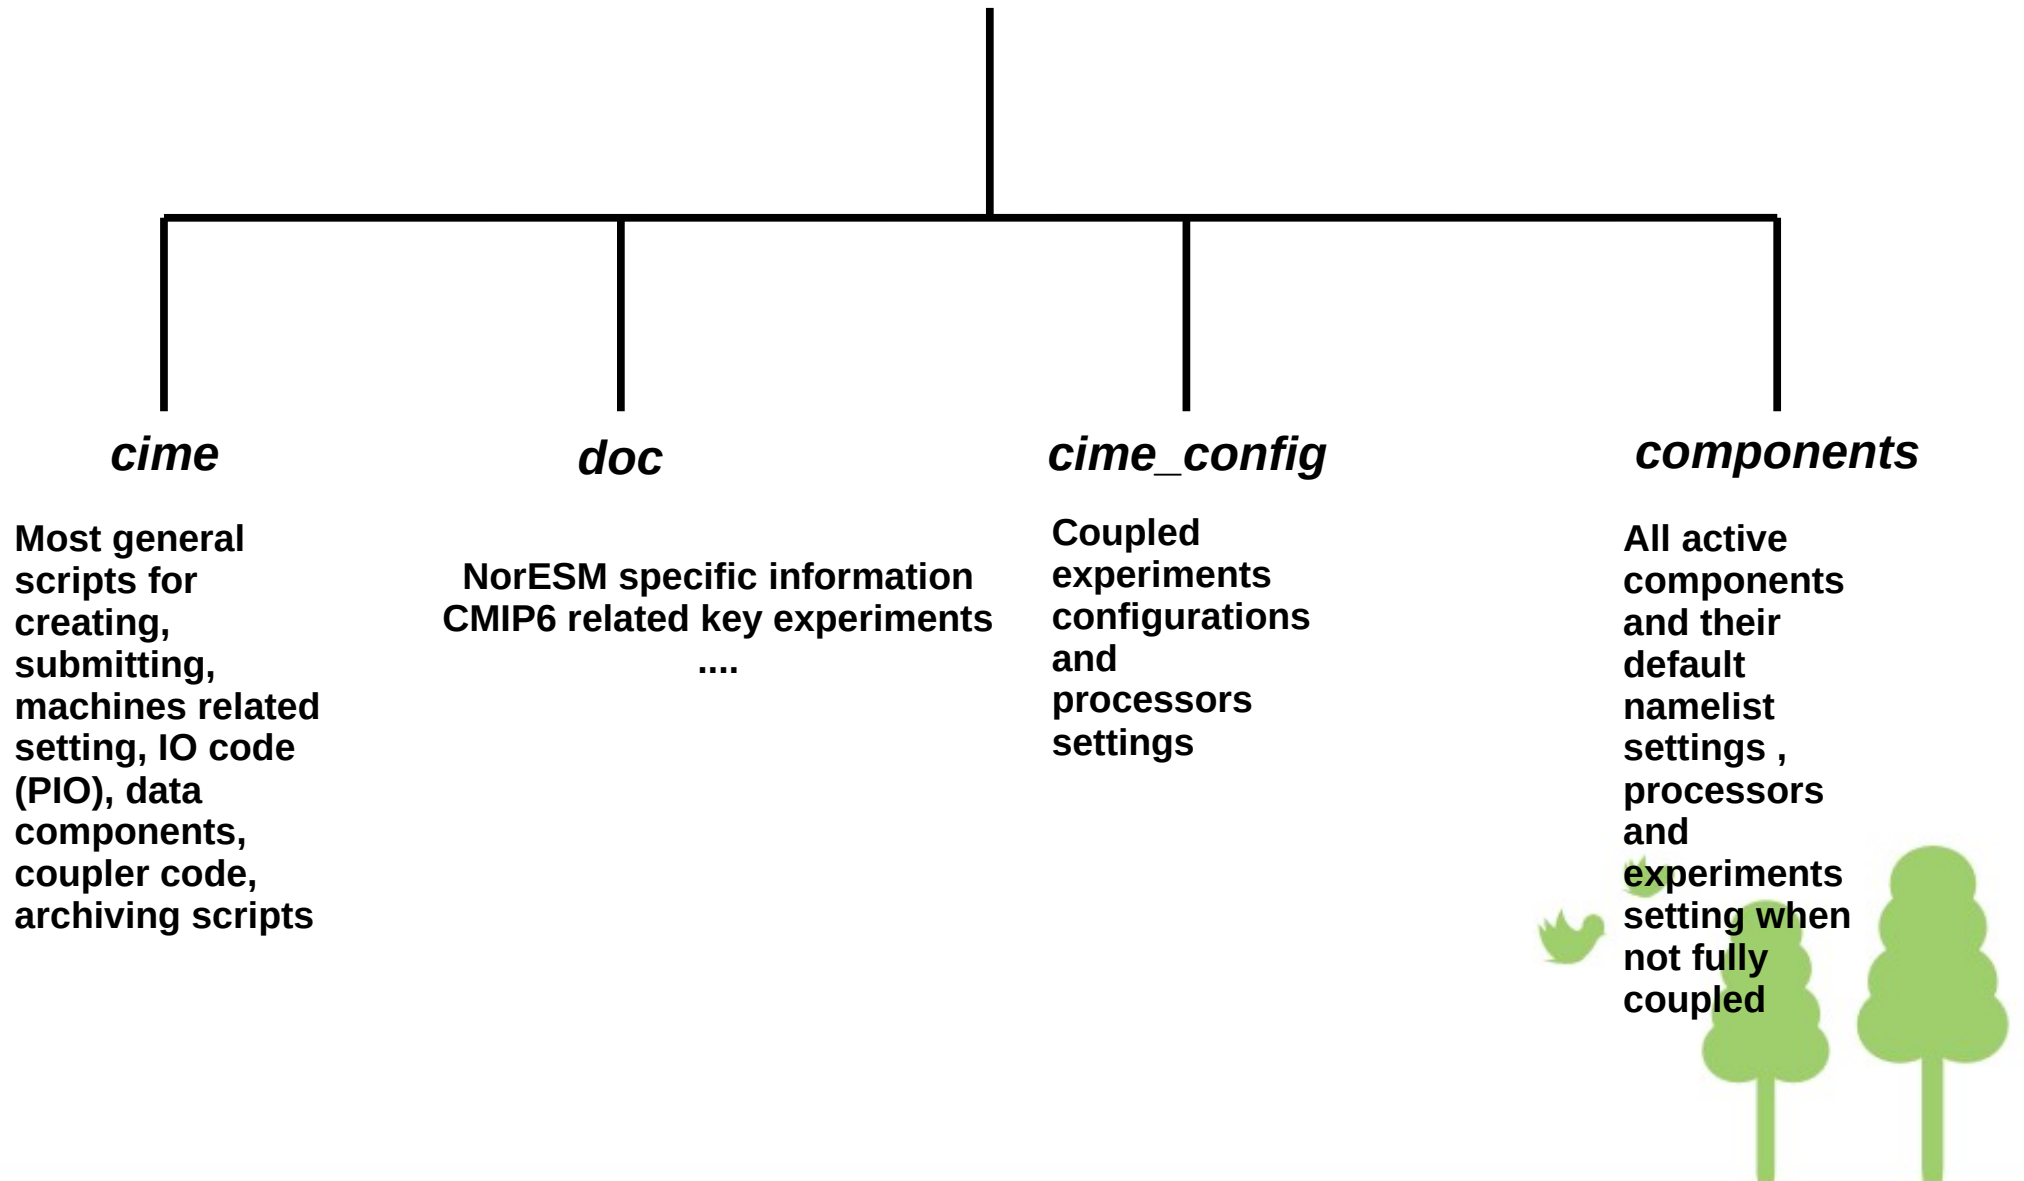
NorESM framework



NorESM/CESM execution strategy



NorESM structure



cime

- ***scripts***

create_newcase, create_clone, query_config

- ***scripts/Tools***: mostly tools for setting, building and executing experiment, checks inputdata, archive data; some of these tools/scripts are linked to experiment directory when you create an experiment
- ***scripts/lib/CIME***: contains mostly python scripts to execute these tools



cime/src/

- **components**

- data_comps** - Data driven components

- stub_comps** - Stub components

- xcpl_comps** - Dead components- only for technical system testing

- **drivers**

- mct** – model coupling toolkit

- **externals**

- pio** – parallel IO code

- **share**

- This module exists to collect code shared between various components

components

- ***cam*** – atmosphere; ***cice*** – sea-ice; ***cism*** – ice sheet; ***clm*** – land ; ***blom***, ***pop*** – ocean; ***mosart***, ***rtm*** – river transport ; ***ww3*** – wave model

Every component have ***src*** and ***cime_config*** directories

- ***cime_config***:

buildlib – build particular component library

buildnml – build particular component namelist

config_pes.xml – processors configuration for particular compsets

config_compsets.xml – compsets with some data components



Steps to execute an experiment

- Create an experiment – *create_newcase*
- Configure an experiment – *case.setup*
- Build an experiment – *case.build*
- Set *WALLTIME* and *SIMULATION RUN TIME* in *env_batch.xml* and *env_run.xml*
- Run an experiment – *case.submit*



Creating an experiment

Create an experiment using **create_newcase** ; usually I create a directory *cases* where I want to keep experiments.

noresm/cime/scripts/create_newcase --help

Will provide you all the options;

```
../noresm/cime/scripts/create_newcase --case N1850OCBDRDDMS_f09_tn14 --  
compset N1850OCBDRDDMS --res f09_tn14 --machine fram --project nn2345k --  
run-unsupported
```

-- case name of experiment

-- compset *Short name of compsets*

-- res *grid resolution*

-- machine name of machine

-- project CPU hours account

--run-unsupported *scientifically unsupported resolution with compset*




Compsets

- An experiment with some sets of components and forcing
- List of all compsets
 - ./query_config --compsets* list all existing compsets;
 - ./query_config --compsets allactive*
 - ./query_config --compsets blom*
- All compsets starting with N are NorESM related compsets



- The compset longname has the specified order atm, lnd, ice, ocn, river, glc wave cesm-options
- The notation for the compset longname is
- ***TIME_ATM[%phys]_LND[%phys]_ICE[%phys]_OCN[%phys]_ROF[%phys]_GLC[%phys]_WAV[%phys]_ESP[%phys]_BGC[%phys]***
- Where for the specific compsets below the following is supported
- TIME = Time period (e.g. 2000, HIST, RCP8...)
- ATM = [CAM40, CAM50, CAM54, CAM60] ; LND = [CLM45, CLM50, SLND] ; ICE = [CICE, DICE, SICE]
- OCN = [DOCN, ,AQUAP, SOCN, BLOM]; ROF = [RTM, MOSART, SROF] ; GLC = [CISM1, CISM2, SGLC]
- WAV = [WW3, DWAV, XWAV, SWAV]; ESP = [SESP] ; BGC = optional BGC scenario
- The OPTIONAL %phys attributes specify submodes of the given system
- For example DOCN%DOM is the data ocean model for DOCN
- ALL the possible %phys choices for each component are listed.
- ALL data models must have a %phys option that corresponds to the data model mode
- Each compset node is associated with the following elements
- - lname ; alias ; support (optional description of the support level for this compset)
- Each compset node can also have the following attributes
- - grid (optional regular expression match for grid to work with the compset)



- **Few coupled compsets:**

TIME_ATM[%phys]_LND[%phys]_ICE[%phys]_OCN[%phys]_ROF[%phys]_GLC[%phys]_WAV[%phys]_ESP[%phys]_BGC[%phys]

- **N1850fr2** : 1850_CAM60%NORESM%FRC2_CLM50%BGC-CROP_CICE%NORESM-CMIP6_BLOM
%ECO_MOSART_SGLC_SWAV_BGC%BDRDDMS
- **NHISTfr2** : HIST_CAM60%NORESM%FRC2_CLM50%BGC-CROP_CICE%NORESM-CMIP6_BLOM
%ECO_MOSART_SGLC_SWAV_BGC%BDRDDMS
- **NSSP126fr2** : SSP126_CAM60%NORESM
%FRC2_CLM50%BGC-CROP_CICE%NORESM-CMIP6_BLOM
%ECO_MOSART_SGLC_SWAV_BGC%BDRDDMS
- **N1850fr2 -short name of compset**



Defining compsets (experiment)/files

- Within **components** (not fully coupled – data atmosphere, stub atmosphere, stub or data ocean, ..)

cam/cime_config/config_compsets.xml

cice/cime_config/config_compsets.xml

clm/cime_config/config_compsets.xml

pop/cime_config/config_compsets.xml

cism/cime_config/config_compsets.xml

blom/cime_config/config_compsets.xml

Fully coupled compsets:

- *cime_config/config_compsets.xml*



Grids:

- `./query_config --grids` provide description of all grids; `--long` provides detailed description.

alias: f19_tn14 (only for compsets that are not _POP)
non-default grids are: atm:1.9x2.5 Ind:1.9x2.5 ocnice:tnx1v4
mask is: tnx1v4

alias: T62_tn14 (only for compsets that are not _CAM)
non-default grids are: atm:T62 Ind:T62 ocnice:tnx1v4

- NorESM supported grid: f19_tn14, f09_tn14, f09_tn0254, T62_tn14, TL319_tn14, TL319_tn0254, f09_f09_mg17, f19_f19_mg17...



Grid resolution

- **f19 1.9x2.5** = 1.9 degree latitude and 2.5 degree longitude resolution = 144x96
 - **f09 0.9x1.25** = 0.9 degree latitude and 1.25 degree longitude resolution = 288x192
 - **f02 0.23x0.31** = 0.23 degree latitude and 0.31 degree longitude resolution = 1152x768
 - **tnx1v1 ,tnx1v3, tnx1v4** = tripolar 1 degree grid= 360x384
 - **tnx0.25v1 ,tnx0.25v3, tnx0.25v4** = tripolar quarter degree ocean grid =1440x1152 ; 1,3,4 version number
- NorESM2 is using mostly version 4 which not includes Caspian sea
- **T62, TL319** = approx. 2 degree Gaussian and 0.5 degree spectral grid used for data atmosphere

cime/config/cesm/config_grids.xml



Mapping files

- There is different resolution for ocean, atmosphere and river runoff grid
- Coupler is used to interpolate/transfer fields from one grid to another grid
- Map files are generated for weight factors (***cime/tools***) and these factors are stored; for that purpose ESMF tool is used.
- Below some list of map files:

map_tnx1v4_to_fv0.9x1.25_aave_da_170609.nc

map_tnx1v4_to_fv1.9x2.5_aave_da_170609.nc

map_fv0.9x1.25_to_tnx1v4_aave_da_170609.nc

map_fv1.9x2.5_to_tnx1v4_aave_da_170609.nc

map_r05_to_tnx1v4_e1000r300_170609.nc

cime/config/cesm/config_grids.xml



Configuring experiment

- `./case.setup` ; Creates namelists and various files and directories needed in order to build and run the case.
- Any changes to *env_mach_pes.xml* and *env_case.xml* must be made before running this.
- To run this initially for the experiment, simply run:

`./case.setup`

- To rerun after making changes to *env_mach_pes.xml*, run:

`./case.setup --reset`

`./case.setup --clean` do not remove `user_nl_*` files and Macros file



Creating executable

- **./case.build**
- You can see all software module in *env_mach_specific.xml* and all compiler flags in *Macros.make*
- Processors configuration in *env_mach_pes.xml*
- *./case.build* – it will create namelist files and compile all the required libraries (mct, gptl, csm_share and pio) and components (cam, blom, clm, cice, ..). Finally, build the binary **cesm.exe**;
found in *\$CIME_OUTPUT_ROOT/\$CASE/bld*
- After this, you can modify only *env_batch.xml*; *env_run.xml* files and *user_nl_** files
- *./case.build --clean* component removes object files of components
- *./case.build --clean-all* removes bld directory



Setting simulation period and restart files option

- Modify *env_run.xml* in experiment directory
- *STOP_OPTION*
nseconds,nsecond,nminutes,nminute,nhours,nhour,nmonths, nyears
- *STOP_N*
numerical value

Writing restart files in middle of simulation:

Restart files are written end of the simulation by default; But, if you are having a long simulation of 100 years; for safety reason you want to write restart files at some frequency you can set below option

- *REST_OPTION*
- nseconds,nsecond,nminutes,nminute,nhours,nhour,nmonths, nyears
- *REST_N* ; numerical value
- *DOUT_S_SAVE_INTERIM_RESTART_FILES*; TRUE or FALSE



Setting wallclock time

- Experiment directory *env_batch.xml*
XML block for *case.run*
in that set *JOB_WALLCLOCK_TIME*
- XML block for *case.st_archive*
in that also set *JOB_WALLCLOCK_TIME*
- Here also you can modify *project* for CPU hours if required
usually, it is set during experiment creation



Executing experiment

- **./case.submit**
- It will first update all the *namelist* if/or made any corrections after in *env_run.xml*, *env_batch.xml* and *user_nl_**
- Then, It will submit two jobs scripts **.case.run** and **case.st_archive**
 - .case.run** – executes the model
 - case.st_archive** – is a dependent job on **.case.run**. On successful completion of **.case.run**, it is executed and copy output to short term archiving
- Before calling it check all namelist parameters and job WALLCLOCK time



Checking inputdata

- `./check_input_data` scripts check missing inputdata during `./case.build` and `./case.submit`
- It downloads missing data from NorESM/CESM server during `./case.submit` by invoking `./check_input_data --download`
- *BLOM* and *CAM-OSLO* related data exists only on *NorESM* server; if you are using older version of *NorESM* then, you have to copy these data.



Information about experiment and its status

- *README.case*

it logs all the information regarding experiment

which grids resolution, compsets, which NorESM repository tag and time step

- *CaseStatus*

it logs all information what you have done with experiment such as setting up experiment; building experiment ; executing experiment with time stamps



Successful completion

- Check file CaseStaus
- -----
- 2019-07-29 04:05:11: case.run starting
- -----
- 2019-07-29 04:05:26: model execution starting
- -----
- 2019-07-31 20:49:23: model execution success
- -----
- **2019-07-31 20:49:24: case.run success**
- -----
- 2019-08-01 10:14:28: st_archive starting
- -----
- 2019-08-01 10:17:12: st_archive success
- -----
- 2019-08-01 10:17:12: case.submit starting
- -----
- 2019-08-01 10:17:13: case.submit success case.run:380071.service2, case.st_archive:380072.service2
- -----
- 2019-08-02 22:00:00: case.run starting
- -----
- 2019-08-02 22:00:16: model execution starting
- -----
- 2019-08-05 02:02:45: model execution success
- -----
- **2019-08-05 02:02:45: case.run error**
- **ERROR: RUN FAIL: Command 'mpiexec_mpt -n 450 /work/agu002/noresm/NOICPL_N1850_f19_tn14**
- **_210619_sal4/bld/cesm.exe >> cesm.log.\$LID 2>&1 ' failed**
- **See log file for details:**
- **/work/agu002/noresm/NOICPL_N1850_f19_tn14_210619_sal4/run/cesm.log.380071.service2.190802-220000**



Timing statistics

- Stored within *timing* sub-directory in experiment directory
- File name `cesm_timing.$CASE.*`
- It provides information on grid type, run length, compset, processors configuration and many others. Most important are timing statics :- Model throughput, Model cost and run time

Overall Metrics:

Model Cost: 303.71 pe-hrs/simulated_year
Model Throughput: 36.67 simulated_years/day

Init Time : 94.799 seconds
Run Time : 235636.536 seconds 6.456 seconds/day
Final Time : 0.132 seconds

Actual Ocn Init Wait Time : 19219.605 seconds
Estimated Ocn Init Run Time : 0.000 seconds
Estimated Run Time Correction : 0.000 seconds

(This correction has been applied to the ocean and total run times)

Runs Time in total seconds, seconds/model-day, and model-years/wall-day

CPL Run Time represents time in CPL pes alone, not including time associated with data exchange with other components

TOT Run Time:	235636.536 seconds	6.456 seconds/mday	36.67 myears/wday
CPL Run Time:	47160.812 seconds	1.292 seconds/mday	183.20 myears/wday
ATM Run Time:	53940.430 seconds	1.478 seconds/mday	160.18 myears/wday
LND Run Time:	0.000 seconds	0.000 seconds/mday	0.00 myears/wday
ICE Run Time:	97041.644 seconds	2.659 seconds/mday	89.03 myears/wday
OCN Run Time:	214741.568 seconds	5.883 seconds/mday	40.23 myears/wday
ROF Run Time:	4463.872 seconds	0.122 seconds/mday	1935.54 myears/wday
GLC Run Time:	0.000 seconds	0.000 seconds/mday	0.00 myears/wday
WAV Run Time:	0.000 seconds	0.000 seconds/mday	0.00 myears/wday
ESP Run Time:	0.000 seconds	0.000 seconds/mday	0.00 myears/wday
CPL COMM Time:	92517.608 seconds	2.535 seconds/mday	93.39 myears/wday

Continuing an experiment

- *CONTINUE_RUN* in *env_run.xml*; TRUE or FALSE
But, you need all restart files and *rpointer.** files in *run* folder
- *RESUBMIT* in *env_run.xml* ; an integer value
it will auto resubmit till specified value; you will have total simulation period $STOP_N * (RESUBMIT + 1)$
- when you are having *WALLCLOCK* time limitation on system. For example, you want to have 200 years simulation and *WALLCLOCK* time limitation is 5 days; you are able to simulate 10 model years/day; to complete 200 model years simulation set *RESUBMIT*=3, *STOP_N* to 50 and *STOP_OPTION* to *nyears*



Changing a namelist parameter

- Within experiment directory there are *user_nl_** files
- If you plan to modify any default namelist parameter and want to use some extra namelist parameter then, you can add them to these relevant files
- For example you want to have extra cam output *user_nl_cam*
`fincl1 = 'FSN200','FSN200C','FLN200'`
- <http://www.cesm.ucar.edu/models/cesm1.2/cesm/doc/usersguide/x2172.html>
- Example for *user_nl_clm*
`finidat =
'/cluster/shared/noresm/inputdata/cesm2_init/b.e20.B1850.f09_g17.pi_control.all.297/0308-01-01/b.e20.B1850.f09_g17.pi_control.all.297.clm2.r.0308-01-01-00000.nc'

use_init_interp = .true.

reset_snow = .true.`
- Best practice for tuning parameter



Changing source code

- Within experiment, there is a directory name **SourceMods**; inside that there are sub-directories like as below; example from one data atmosphere experiment

***src.cice src.datm src.drof src.drvt src.blom src.sglc
src.share src.slnd src.swav***

- BLOM will use f77 style header files (*.h) and if modified, then one must put all dependent source files to **SourceMods/src.blom**
- put relevant source file to relevant model sub-directory and build the model again using ./case.build and then, submit
- Once satisfied, you can commit it to repository
- Best practice for modifying code



Namelist and SourceMods

- *--user-mods-dir* option in *create_newcase*

--user-mods-dir USER_MODS_DIR

USER_MODS_DIR name of directory which contains *user_nl_** and *SourceMods*

- If used it will copy whole structure to experiment.
- Usually, used when you want some typical settings for particular experiments.
- Exists under *cime_config/usermods_dirs*
components/cime_config/usermods_dirs



xmlquery and xmlchange

- xmlquery :- provides the information and it is value which are sets in *.xml files
- xmlchange :- is used to change of values/parameters set in *.xml files
- xmlquery --help xmlchange --help
- Examples:
 - `./xmlchange --id STOP_OPTION --val nyears`
 - `./xmlquery --value STOP_OPTION`
 - `./xmlchange STOP_OPTION=nyears,STOP_N=1`



Branch run

- <http://www.cesm.ucar.edu/models/cesm1.2/cesm/doc/usersguide/x1894.html>
- Mostly used for tuning experiments and investigating parameter affects
- *RUN_TYPE* to “*branch*”
- *RUN_REFDIR* directory containing reference data
- *RUN_REFCASE* name of reference case
- *RUN_REFDATE* Reference date branch run
- *GET_REFCASE TRUE* else you have to copy data to run folder



Hybrid run

- It could have reference files from many experiments
- *RUN_TYPE* to “*hybrid*”
- *RUN_REFDIR* directory containing reference data
- *RUN_REFCASE* name of reference case
- *RUN_REFDATE* Reference date for hybrid run
- *GET_REFCASE TRUE* else you have to copy data to run folder
- *RUN_STARTDATE* Run start date (yyyy-mm-dd).
- Can be used for together combinations of initial/restart files
- In a hybrid initialization, the ocean model does not start until the second ocean coupling



Cloning run

- It creates an exact copy of the experiment
- ***create_clone --case abcdef --clone old_case***
- You can do desired modifications before building and submitting executable
- You should not modify *env_case.xml*, *env_build.xml*
- *README.case* will be having same information as *old_case*; so edit it if have made some modifications or used different model version to backtrace



Setting processors count example for fully coupled code

```
<grid name="a%1.9x2.5.+l%1.9x2.5.+oi%tnx1v4" >
  <mach name="hexagon|vilje|fram">
    <pes pesize="any" compset="CAM60%PTAERO.+CLM50%BGC-CROP.+CICE.+MCOM%ECO">
      <comment>none</comment>
      <ntasks>
        <ntasks_atm>768</ntasks_atm>
        <ntasks_lnd>256</ntasks_lnd>
        <ntasks_rof>8</ntasks_rof>
        <ntasks_ice>504</ntasks_ice>
        <ntasks_ocn>186</ntasks_ocn>
        <ntasks_glc>768</ntasks_glc>
        <ntasks_wav>300</ntasks_wav>
        <ntasks_cpl>768</ntasks_cpl>
      </ntasks>
      <nthrds>
        <nthrds_atm>1</nthrds_atm>
        <nthrds_lnd>1</nthrds_lnd>
        <nthrds_rof>1</nthrds_rof>
        <nthrds_ice>1</nthrds_ice>
        <nthrds_ocn>1</nthrds_ocn>
        <nthrds_glc>1</nthrds_glc>
        <nthrds_wav>1</nthrds_wav>
        <nthrds_cpl>1</nthrds_cpl>
      </nthrds>
      <rootpe>
        <rootpe_atm>0</rootpe_atm>
        <rootpe_lnd>0</rootpe_lnd>
        <rootpe_rof>256</rootpe_rof>
        <rootpe_ice>264</rootpe_ice>
        <rootpe_ocn>768</rootpe_ocn>
        <rootpe_glc>0</rootpe_glc>
        <rootpe_wav>0</rootpe_wav>
        <rootpe_cpl>0</rootpe_cpl>
      </rootpe>
    </pes>
  </mach>
</grid>
```

Total number of processors =
ntasks_atm + ntasks_ocn

0 ntasks_atm

ATM			OCN
LND	ROF	ICE	
WAV			
GLC			
CPL			

rootpe_ocn=ntasks_atm

rootpe_rof=ntasks_lnd

rootpe_ice=ntasks_lnd+ntasks_rof



Processors count

- *env_mach_pes.xml*; within experiment
- *config_pes.xml*; within **cime_config** or **components/*/cime_config**
- For *atmosphere*, you can set maximum number of processors *nlat/3*atmosphere levels*; for NorESM2 32 atmosphere levels.....
- For *blom*, you have to look relevant directory
- **components/blom/bld/"gridname"**
there are *patch.input.** files, You should use last numerical values of these files as a processors count
- *Land* is quite trivial, you can set whatever number you want
- *CICE* uses auto decomposition and you can set any processors count and based on that it will choose parallel method



Configuration files for porting to new machines

- ***cime/config/cesm***

config_machines.xml - machine related settings, where are inputdata or where to execute model or copy output, which modules to use...

config_batch.xml - setting batch queue setting (slurm, PBS, ..)

config_compilers.xml - compilers related settings

- ***cime_config/components/*/cime_config***

config_pes.xml – processors distribution settings



```

<machine MACH="fram">
  <DESC>Lenovo NeXtScale H5, 32-way nodes, dual 10-core Xeon E5-2683@2.10GHz, 64 GiB per node, os is Linux, batch system is SLURM</DESC>
  <OS>LINUX</OS>
  <COMPILERS>intel</COMPILERS>
  <MPILIBS>impi</MPILIBS>
  <CTMF_OUTPUT_ROOT>/cluster/work/users/SHSER/ncores</CTMF_OUTPUT_ROOT>
  <modules>
    <command name="purge">--force</command>
    <command name="load">StdEnv</command>
    <!-- djlo Deactivated THT settings -->
    <!-- command name="load">intel/2016a</command-->
    <!-- command name="load">netCDF-Fortran/4.4.3-intel-2016a</command-->
    <!-- command name="load">PnetCDF/1.8.1-intel-2016a</command-->
    <!-- command name="load">CMake/3.5.2-intel-2016a</command-->
    <command name="load">intel/2018a</command>
    <command name="load">netCDF-Fortran/4.4.4-intel-2018a-HDF5-1.8.19</command>
    <command name="load">PnetCDF/1.8.1-intel-2018a</command>
    <command name="load">CMake/3.9.1</command>
  </modules>
</module_system>
<environment_variables>
  <env name="KMP_STACKSIZE">64M</env>
  <env name="I_MPI_EXTRA_FILESYSTEM_LIST">lustre</env>
  <env name="I_MPI_EXTRA_FILESYSTEM">on</env>
</environment_variables>
<resource_limits>
  <resource name="RLIMIT_STACK">-1</resource>
</resource_limits>

  <command name="load">netCDF-Fortran/4.4.4-intel-2018a-HDF5-1.8.19</command>
  <command name="load">PnetCDF/1.8.1-intel-2018a</command>
  <command name="load">CMake/3.9.1</command>
</modules>
</module_system>
<environment_variables>
  <env name="KMP_STACKSIZE">64M</env>
  <env name="I_MPI_EXTRA_FILESYSTEM_LIST">lustre</env>
  <env name="I_MPI_EXTRA_FILESYSTEM">on</env>
</environment_variables>
<resource_limits>
  <resource name="RLIMIT_STACK">-1</resource>
</resource_limits>
</machine>

```

config_batch.xml

```
<batch_system MACH="fram" type="slurm">
  <batch_submit>sbatch</batch_submit>
  <submit_args>
    <arg flag="--time" name="$JOB_WALLCLOCK_TIME"/>
    <arg flag="-p" name="$JOB_QUEUE"/>
    <arg flag="--account" name="$PROJECT"/>
  </submit_args>
  <directives>
    <directive> --ntasks={{ total_tasks }}</directive>
    <directive> --export=ALL</directive>
    <directive> --switches=1</directive>
  </directives>
  <queues>
    <queue walltimemax="00:59:00" nodemin="1" nodemax="288" default="true">normal</queue>
  </queues>
</batch_system>
```



config_compilers.xml

```
<compiler MACH="fram">
  <CPPDEFS>
    <append> -D$(OS) </append>
  </CPPDEFS>
  <FFLAGS>
    <append> -xCORE-AVX2 -no-fma </append>
  </FFLAGS>
  <NETCDF_PATH>$(EBROOTNETCDFMINFORTRAN)</NETCDF_PATH>
  <PNETCDF_PATH>$(EBROOTPNETCDF)</PNETCDF_PATH>
  <MPI_PATH>$(MPI_ROOT)</MPI_PATH>
  <MPI_LIB_NAME>mpi</MPI_LIB_NAME>
  <FFLAGS>
    <append DEBUG="FALSE"> -O2 </append>
    <append MODEL="blom"> -r8 </append>
  </FFLAGS>
  <MPICC> mpiicc </MPICC>
  <MPICXX> mpiicpc </MPICXX>
  <MPIFC> mpiifort </MPIFC>
  <PIO_FILESYSTEM_HINTS>lustre</PIO_FILESYSTEM_HINTS>
  <SLIBS>
    <append>-mkl=sequential -lnetcdff -lnetcdf</append>
  </SLIBS>
</compiler>
```

Set compilers FLAGS
NETCDF and PNETCDF path
Compilers/wrapper used on machine
MODEL related flags

